

```
import pandas as pd
import numpy as np
```

```
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import KBinsDiscretizer
from sklearn.compose import ColumnTransformer
```

```
!wget -O titanic_train.csv https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv
```

```
--2023-06-25 06:29:42-- https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.109.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 60302 (59K) [text/plain]
Saving to: 'titanic_train.csv'
```

```
titanic_train.csv 100%[=====>] 58.89K --.-KB/s in 0.008s
```

```
2023-06-25 06:29:42 (7.30 MB/s) - 'titanic_train.csv' saved [60302/60302]
```

```
df=pd.read_csv('titanic_train.csv',usecols=['Age','Fare','Survived'])
```

```
df.head()
```

	Survived	Age	Fare
0	0	22.0	7.2500
1	1	38.0	71.2833
2	1	26.0	7.9250
3	1	35.0	53.1000
4	0	35.0	8.0500

```
df.dropna(inplace=True)
```

```
df.shape
```

```
(714, 3)
```

```
x=df.iloc[:,1:]
```

```
y=df.iloc[:,0]
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=.2,random_state=42)
```

```
y_train.head()
```

```
328    1
73     0
253    0
719    0
666    0
Name: Survived, dtype: int64
```

```
clf=DecisionTreeClassifier()
```

```
clf.fit(x_train,y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier()
```

```
y_pred=clf.predict(x_test)
```

```
accuracy=accuracy_score(y_test,y_pred)
```

```
accuracy
```

0.6293706293706294

```
kbin_age=KBinsDiscretizer(n_bins=25,encode='onehot',strategy="quantile")
kbin_fare=KBinsDiscretizer(n_bins=25,encode='onehot',strategy="quantile")
```

```
trf=ColumnTransformer([('first',kbin_age,[0]),('second',kbin_fare,[0])])
```

```
x_train_trf=trf.fit_transform(x_train)
x_test_trf=trf.transform(x_test)
```

```
trf.named_transformers_['first'].bin_edges_
```

```
array([array([ 0.42,  3.   ,  8.6 , 16.   , 18.   , 19.   , 20.9 , 22.   , 23.   ,
              24.   , 25.   , 27.   , 28.   , 29.   , 30.   , 32.   , 33.   , 35.   ,
              36.2 , 39.   , 42.   , 45.   , 48.   , 51.4 , 59.2 , 80.   ])],
      dtype=object)
```

```
output=pd.DataFrame({'age':x_train['Age'],'age_trf':x_train_trf[:,0],'fare':x_train['Fare'],'fare_trf':x_train_trf[:,1]})
```

```
output['age_labels']=pd.cut(x=x_train['Age'],bins=trf.named_transformers_['first'].bin_edges_[0].tolist())
output['fare_labels']=pd.cut(x=x_train['Fare'],bins=trf.named_transformers_['second'].bin_edges_[0].tolist())
```

```
output.sample(5)
```

	age	age_trf	fare	fare_trf	age_labels	fare_labels
607	27.0		30.5000		(25.0, 27.0]	(30.0, 32.0]
103	33.0		8.6542		(32.0, 33.0]	(8.6, 16.0]
462	47.0		38.5000		(45.0, 48.0]	(36.2, 39.0]
765	51.0		77.9583		(48.0, 51.4]	(59.2, 80.0]
869	4.0		11.1333	(0, 0]	(3.0, 8.6]	(8.6, 16.0]

```
clf=DecisionTreeClassifier()
clf.fit(x_train_trf,y_train)
y_pred2=clf.predict(x_test_trf)
```

```
accuracy_score(y_test,y_pred2)
```

0.6013986013986014

Write a classification program for implementing SVM using MNIST dataset

```
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

```
digits = datasets.load_digits()
```

```
X_train, X_test, y_train, y_test=train_test_split(digits.data, digits.target, test_size=0.2, random_state=42)
```

```
svm_classifier = SVC()
```

```
svm_classifier.fit(X_train, y_train)
```

```
svm_classifier
```

```
y_pred = svm_classifier.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 0.9861111111111112

Write a classification program for implementing Naïve Bayes algorithm using iris dataset

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

iris=load_iris()
X_train, X_test, y_train, y_test=train_test_split(iris.data,iris.target,test_size=0.2,random_state=42)

nb_classifier = GaussianNB()

nb_classifier.fit(X_train, y_train)



▾ GaussianNB  

    GaussianNB()



y_pred = nb_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

Accuracy: 1.0
```

Write a classification program for implementing decision tree using pima-indians-diabetes dataset

```
from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score

diabetes = load_diabetes()
X_train, X_test, y_train, y_test = train_test_split(diabetes.data, diabetes.target, test_size=0.2, random_state=42)
dt_classifier = DecisionTreeClassifier()
dt_classifier.fit(X_train, y_train)
y_pred = dt_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

Accuracy: 0.0
```

Write a classification program for implementing kNN

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

iris = load_iris()
X_train, X_test, y_train, y_test = train_test_split(iris.data, iris.target, test_size=0.33, random_state=42)
knn_classifier = KNeighborsClassifier(n_neighbors=10)
knn_classifier.fit(X_train, y_train)
y_pred = knn_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

Accuracy: 0.98
```

Write a clustering program for implementing k Means , k-medoids and Hierarchical Clustering using Wisconsin Breast Cancer Dataset

```
!pip install scikit-learn-extra
```

```
Collecting scikit-learn-extra
  Downloading scikit_learn_extra-0.3.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (2.0 MB)
    2.0/2.0 MB 23.1 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib/python3.10/dist-packages (from scikit-learn-extra) (1.22.4)
Requirement already satisfied: scipy>=0.19.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn-extra) (1.10.1)
Requirement already satisfied: scikit-learn>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn-extra) (1.2.2)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.23.0->scikit-learn-extra) (1.2.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.23.0->scikit-learn-extra) (3.1.0)
Installing collected packages: scikit-learn-extra
Successfully installed scikit-learn-extra-0.3.0
```

◀ ▶

4/6

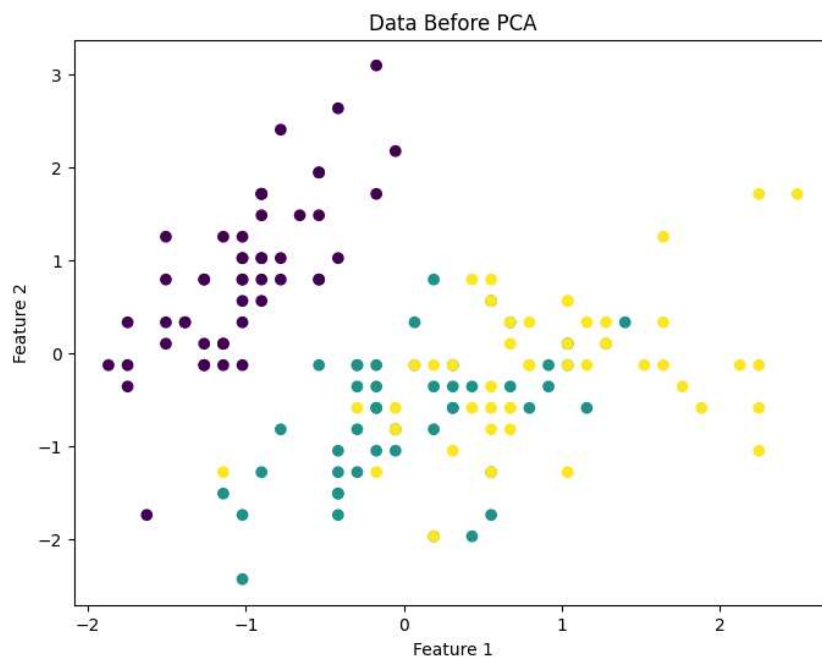
```
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

iris = load_iris()


scaler = StandardScaler()
X = scaler.fit_transform(iris.data)

pca = PCA(n_components=2) # Specify the number of components to keep
X_pca = pca.fit_transform(X)

# Visualize the data before PCA
plt.figure(figsize=(8, 6))
plt.scatter(X[:, 0], X[:, 1], c=iris.target)
plt.title("Data Before PCA")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.show()
```



```
plt.figure(figsize=(8, 6))
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=iris.target)
plt.title("Data After PCA")
plt.xlabel("Principal Component 1")
plt.ylabel("Principal Component 2")
plt.show()
```



Write a program to evaluate Classification Model using different Evaluation Metrics

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report

iris = load_iris()
X_train, X_test, y_train, y_test = train_test_split(iris.data, iris.target, test_size=0.33, random_state=42)

knn_classifier = KNeighborsClassifier(n_neighbors=10)
knn_classifier.fit(X_train, y_train)
y_pred = knn_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, average='weighted')
recall = recall_score(y_test, y_pred, average='weighted')
f1 = f1_score(y_test, y_pred, average='weighted')
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.98
Precision: 0.98125
Recall: 0.98
F1 Score: 0.98
Classification Report:
              precision    recall  f1-score   support

    0         1.00        1.00        1.00        19
    1         0.94        1.00        0.97        15
    2         1.00        0.94        0.97        16

 accuracy          0.98          0.98          0.98          50
  macro avg         0.98          0.98          0.98          50
 weighted avg         0.98          0.98          0.98          50
```